

CIS 520, *Operating Systems Concepts*

Lecture 8 *Input/Output*



I/O Devices

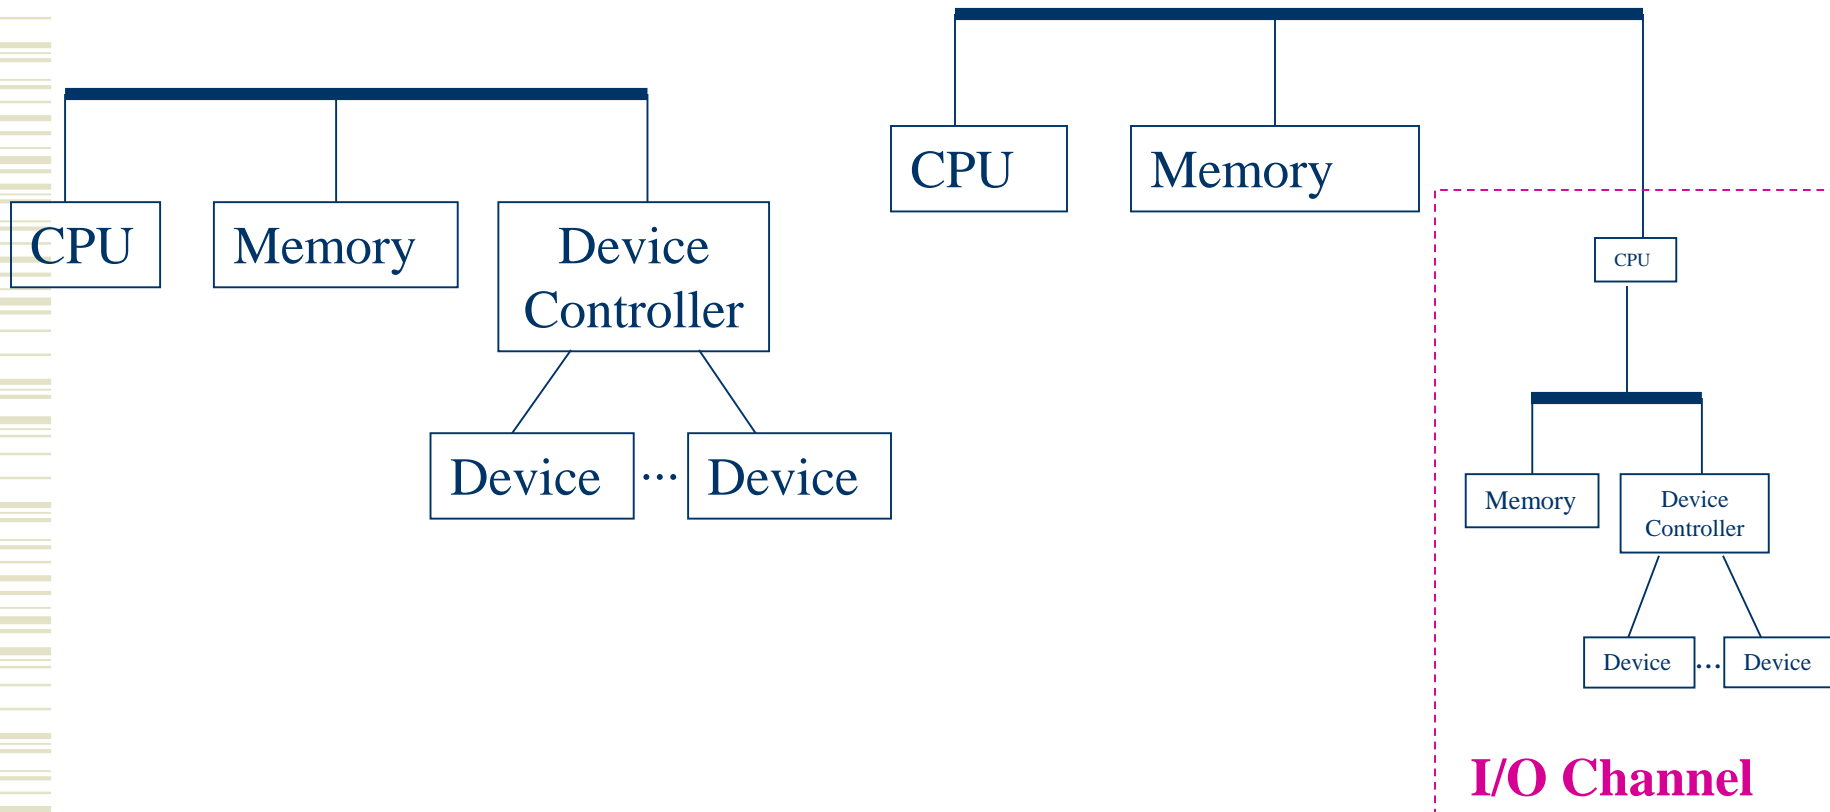
- ◆ There are two categories of the I/O devices:
 - *Block Devices* store and deliver information in blocks of bytes
 - *Disk* uses fixed-size blocks
 - *Tape* and *asynchronous terminals* use variable size blocks
 - *Character Devices* (*dumb terminals, keypunch, card readers, printers, MIDI*) store and deliver information as a stream of bytes

And then there are devices (e.g., *computer clocks, telemetric devices*, etc.) that don't belong to these categories.

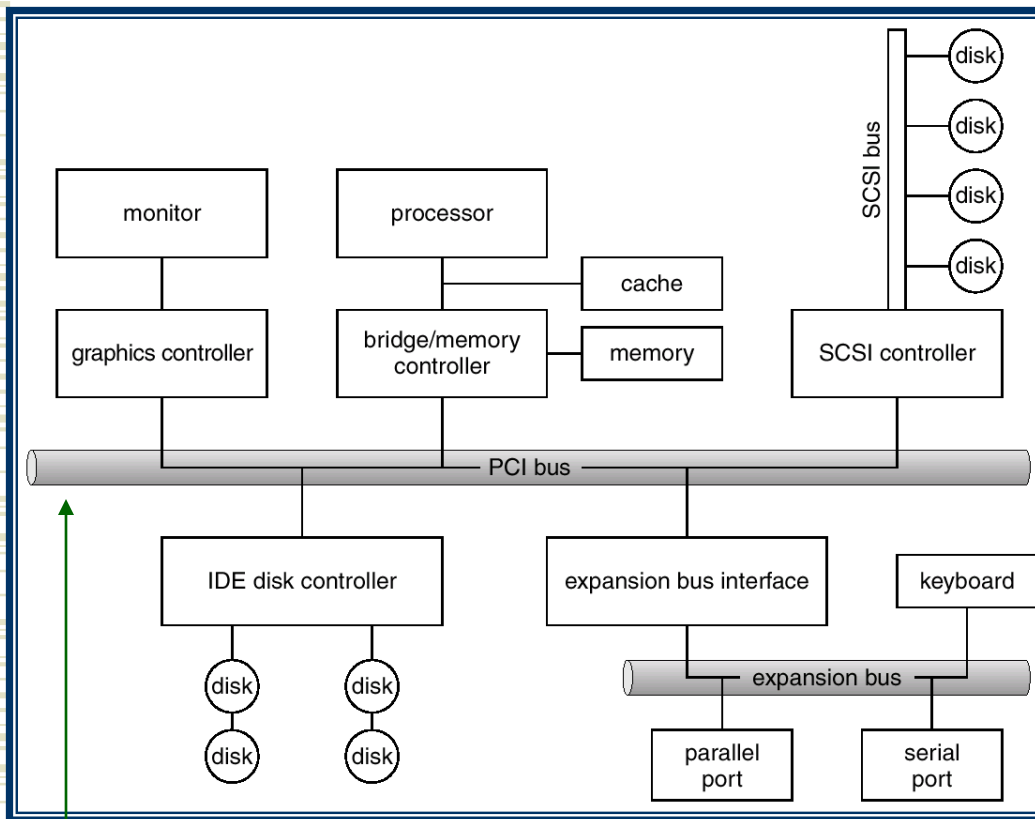
Device Controllers and I/O Channels

- ◆ The *programmable* component of a device is called a *controller*. Controllers connect to the system bus
- ◆ In some cases, a separate CPU is assigned to deal with a device (or a set of devices). Such a CPU and its supporting circuitry form an *I/O Channel*

Device Controllers and I/O Channels (cont.)



A Real-Life Example (from the Book)

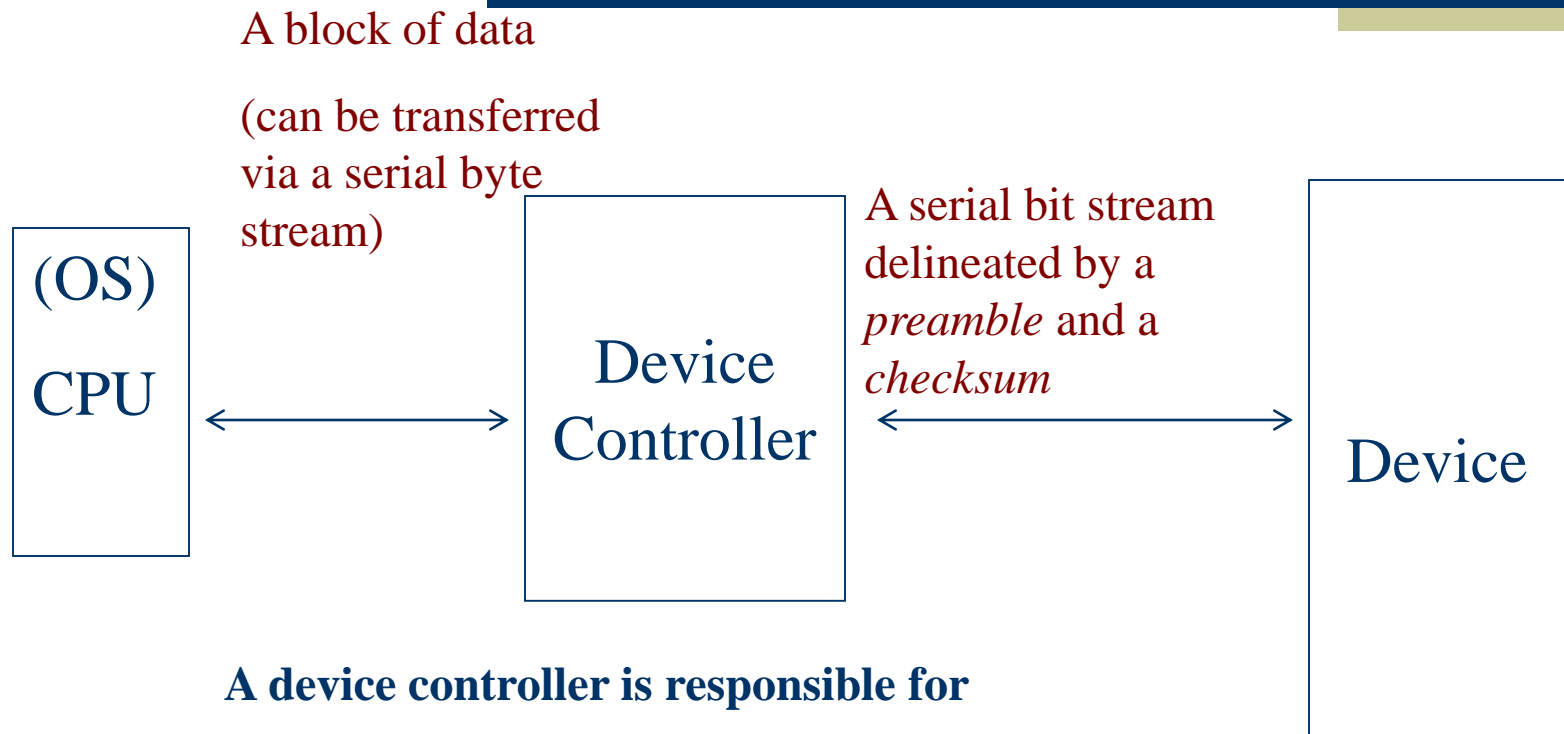


Note: Integrated Disk Electronics (IDE) is the least expensive current disk technology. IDE support is usually built into the main board.

Small Computer Systems Interface (SCSI) supports disks, tapes, and CD-ROM drives. While IDE disks provide up to one gigabyte of storage, SCSI disks are available with four to 9 gigabytes of storage.

Peripheral Component Interconnection (PCI)

Interfaces and Responsibilities




A device controller is responsible for

1. Ensuring that the I/O is error-free and
2. Translating the OS requests into a series of low-level device commands and vice versa

Memory Access

- ◆ There are two aspects to that:
 1. The *control path* aspect—sending actual commands (like *read* or *write*) from the OS to the controller. Typically, this is achieved by writing to the controller registers, which can also be *memory-mapped*
 2. The *data transfer path* aspect—moving data between the memory and the controller is typically achieved via *Direct Memory Access (DMA)*

Disk Hardware




Cylinder = set of all *tracks* under the heads in a present position

Commands:

- *SEEK cylinder*
- *READ track, [sector]*
- *WRITE track, [sector]*
- *RECALIBRATE*

Magnetic Tape Hardware



Commands:

- ***READ record***
- ***WRITE record / EOF***
- ***ERASE GAP***
- ***REWIND BOT / EOF***

The Control Path

- ◆ The OS (more specifically, a device driver) communicates with a controller by writing commands and their parameters in the *controller* registers. For example, a disk controller might have three registers called *command*, *address*, and *start*, so a request to *read* sector 3 of track 5 of cylinder 7 into a memory location (h)AF67C018 would result in setting the values of the first two registers as follows:

Command Register:



| | | | |
|---------|---------|---------|-----------|
| 0 0 0 1 | 0 0 1 1 | 0 1 0 1 | 0 0 1 1 1 |
|---------|---------|---------|-----------|

Address Register:

| | | | |
|---------|---------|---------|---------|
| 1 0 1 0 | 1 1 1 1 | 0 1 1 0 | 0 1 1 1 |
| 1 1 0 0 | 0 0 0 0 | 0 0 0 1 | 1 0 0 0 |

and then writing anything (say 0) to the *start* register

The Control Path (cont.)

- ◆ In addition, a *status* register would display the result of the operation after it was finished
- ◆ The *status* register could also be used for *polled I/O*, but a typical way to communicate with the disk is via an interrupt

An Example: Memory-Mapped I/O in IBM PC

| I/O Controller | I/O Address | Interrupt Vector |
|--------------------|-------------|------------------|
| Clock | 040-043 | 8 |
| Keyboard | 060-063 | 9 |
| Secondary RS232 | 2F8-2FF | 11 |
| Hard Disk | 320-32F | 13 |
| Printer | 378-37F | 15 |
| Monochrome Display | 380-3BF | - |
| Color Display | 3D0-3DF | - |
| Diskette | 3F0-3F7 | 14 |
| Primary RS232 | 3F8-3FF | 12 |

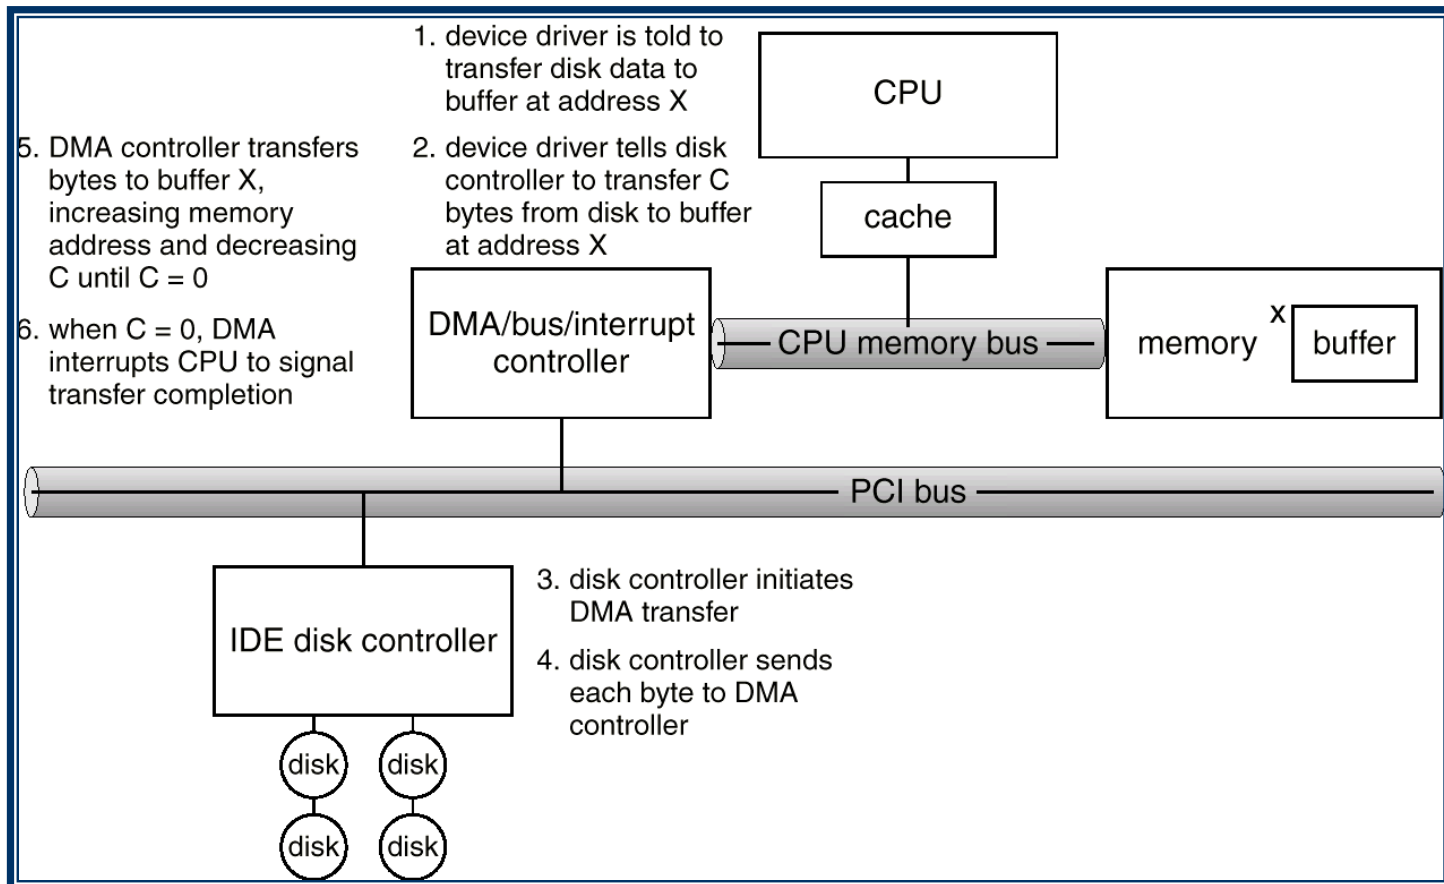
The Data Path

- ◆ The controller stores the data (read from or to be written to the device) in its internal buffer—or it would not be able to keep up with the transfer rate
- ◆ These data need to be transferred to or from the main memory
- ◆ The OS could deal with the controller by reading or writing its data, but interrupting CPU for each byte or at most word of data in a large transfer is very expensive

Direct Memory Access (DMA)

- ◆ DMA solves the problem by transferring the data between the main memory and the controller's buffer *without* getting the CPU (and OS) involved beyond ordering a transfer
- ◆ A DMA (which could be a separate device or part of the device controller) accesses main memory by *stealing* cycles from CPU


DMA Operation (from the Book)




Interleaving Sectors

- ◆ While the DMA transfer of a sector's data takes place, a controller may be unable to keep up with reading the next sector (or even two sectors)
- ◆ Then, the full disk rotation must take place before the next sector is read
- ◆ This could substantially decrease performance, so the disk should be formatted so as to mitigate this problem by skipping adjacent sectors
- ◆ Such a practice is called *interleaving*


Interleaving Sectors (Example)



No Interleaving



Single Interleaving



Double Interleaving

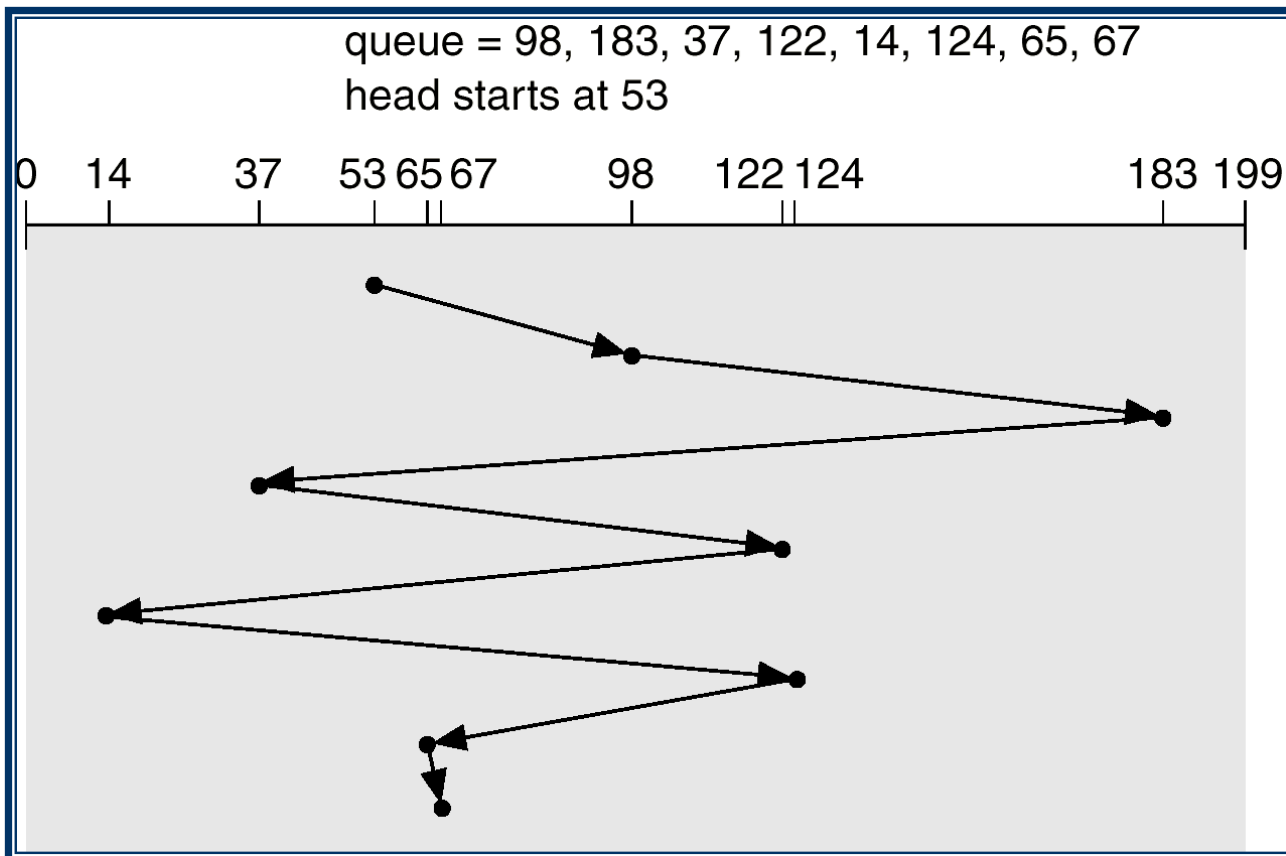
By the way, disk tracks typically have 8 to 32 sectors. The number of sectors is a constant for the whole disk.

Question: Do we waste much space because of the difference in the track size?


Scheduling Disk Requests To Optimize Delay

- ◆ *First-Come First-Served* algorithm optimizes nothing
- ◆ *Shortest Seek Time First (SSTF)* algorithm optimizes seek time, but tends to keep the head around the center, starving the processes that request outer tracks
- ◆ *Elevator (SCAN)* algorithm and its derivatives guarantee the upper bound on the total motion


First-Come First-Served



Shortest Seek Time First (SSTF)



Elevator (SCAN)



A Bottleneck

- The performance of CPUs and RAM has been increasing exponentially (in time)
- The performance of the disk drives has been increasing *sub-linearly*
- In 1987, UC Berkeley researchers David Patterson, Garth Gibson, and Randy Katz have invented the *RAID* technology to deal with the problem

Redundant Array of Inexpensive (or Independent) Disks (RAID)

- ♦ *RAID* is the technology supporting storage systems' resilience to disk failure through the use of multiple disks and by the use of data distribution and correction techniques
- ♦ A *RAID array* is a collection of drives that collectively act as a single storage system, which can tolerate the failure of a drive without losing data, and which can operate independently of one another
- ♦ Different *levels* of RAID are defined according to how the data and parity information is distributed among the disks
- ♦ Parity information supports error correction
(Example: With 38 disks in the system, using bits 1, 2, 4, 8, 16, and 32 for Hamming code, the system can recover from a single disk failure)

Clocks

There are two types:

- Simpler clocks follow the AC frequency of the power line and cause an interrupt at 50-60 Hz
- More useful clocks use a crystal oscillator, which generates signals of high accuracy in the range of 5 to 100 MHz
 - These signals are fed into the *counter*, which counts down to zero
 - At zero count, the counter causes an interrupt (a *tick*)
 - Also at zero count (and initially), the counter is set to the value of the settable *holding register*. In *one-shot* mode, that happens automatically; in *square-wave* mode, the clock stops at the interrupt and waits for a reset

Clock Drivers

- ◆ The *clock driver* is responsible for
 - Maintaining the time of day
 - Preempting processes
 - Accounting for resource use
 - Handling *Alarm* calls
 - Profiling, monitoring, and gathering statistics

Terminals

- ◆ There are three basic types:
 1. Memory-mapped terminals
 2. Serial (synchronous) terminals connected over the RS-232 interface
 3. Asynchronous terminals (although they have been pretty much replaced by PCs now), which support messages of various length
- ◆ *Keyboard* and *display* are two separate devices, and are treated as such by the operating system



Musical Instrument Digital Interface (MIDI) Devices

- ◆ In August of 1983, music manufacturers agreed on a document that is called "MIDI 1.0 Specification".
- ◆ The MIDI protocol supports interworking of keyboards, synthesizers, and sequencers built by different manufacturers
- ◆ It is important to remember that MIDI transmits commands--not audio signals

Physical Interface

- ◆ MIDI uses serial interface. The *serial interface* was chosen by MIDI manufacturers because it is less expensive than parallel interface.
- ◆ The speed of a MIDI serial interface is 31,250 bits per second.
- ◆ There are 10 bits needed for every MIDI digital word or 3125 messages per second.

MIDI Data examples

- ◆ Note on/off (on a particular channel)
- ◆ Additional effects (like sustain pedal)
- ◆ Dynamic range of the note



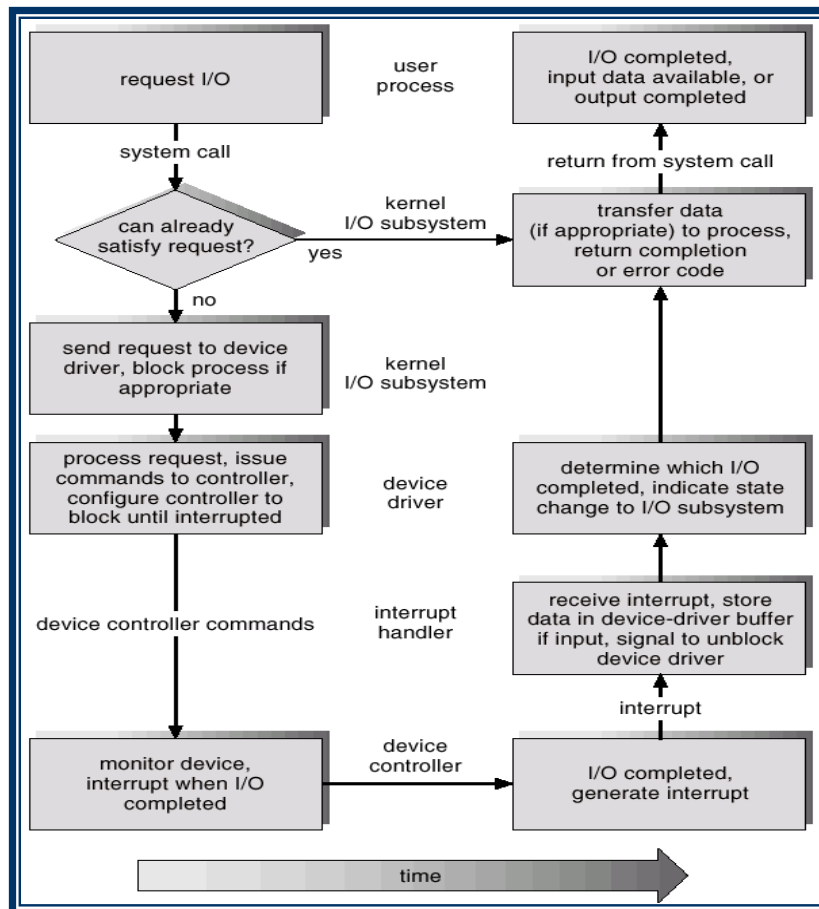
Principles of the I/O Software Design

1. Device independence

- At the user-level, the API is as abstract as possible (e.g., *sort <input >output*)
- All device-specific code goes into *device drivers*

2. Uniform Naming (whence *Universal Resource Identifier [URI]* came)

Life Cycle of an I/O Request



from the Book